

Handling different input types in FastAPI

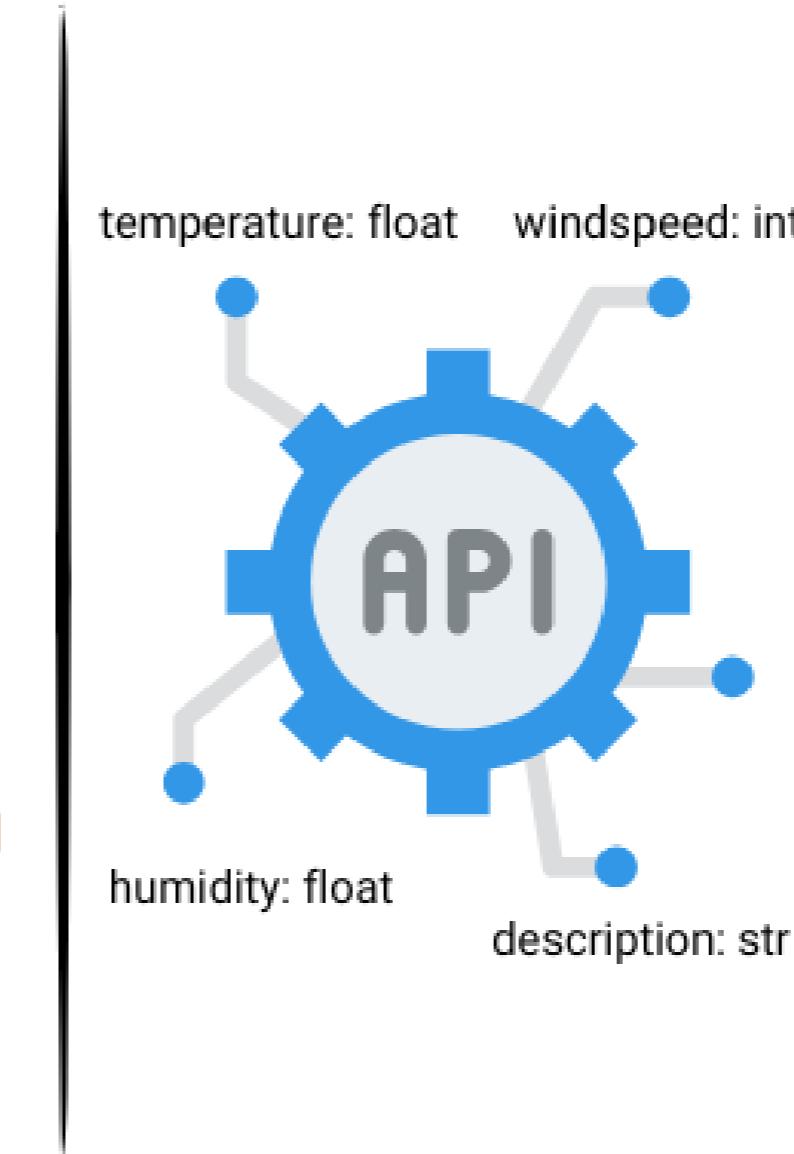
DEPLOYING AI INTO PRODUCTION WITH FASTAPI



Matt Eckerle

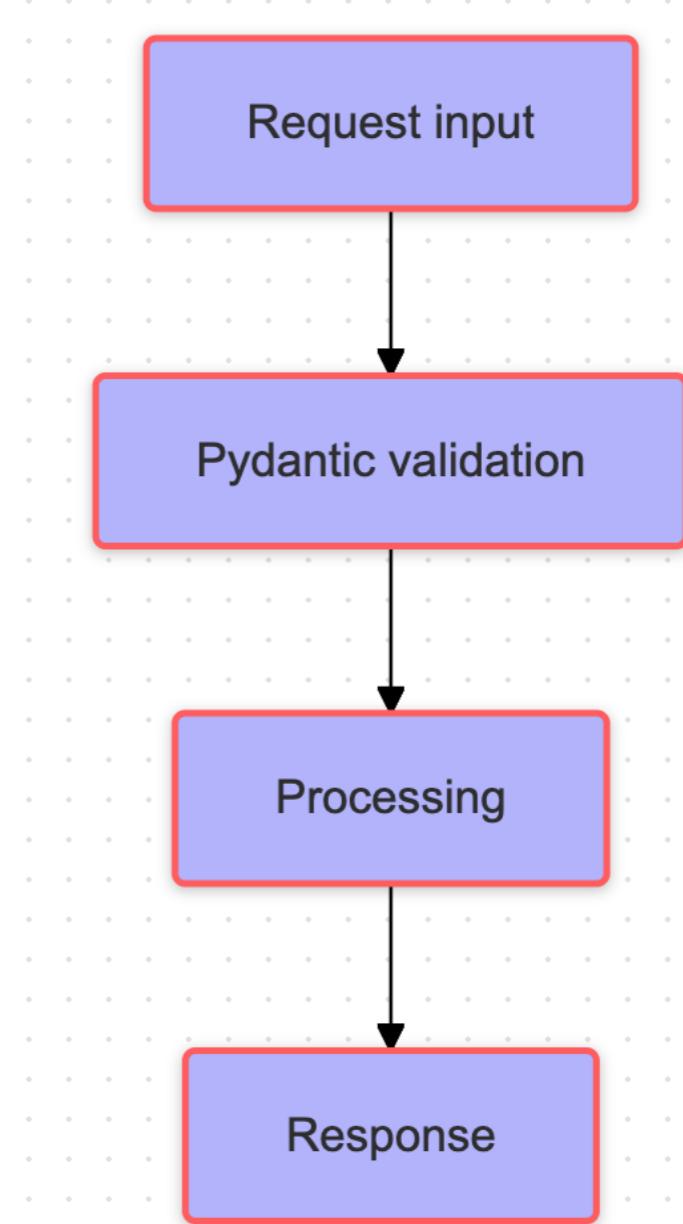
Software and Data Engineering Leader

Restaurant vs API



Validation flow

- Incoming data via request
- Input data validation happens using Pydantic
- Process different types of data as per model requirements
- Processed input sent to the model



Comment moderation system

```
class CommentMetrics(BaseModel):  
    length: int  
    user_karma: int  
    report_count: int  
  
class CommentText(BaseModel):  
    content: str
```



Endpoint for floating point numbers

```
app = FastAPI()
@app.post("/predict")
def predict_score(data: CommentMetrics):
    features = np.array([
        data.length,
        data.user_karma,
        data.report_count
    ])
    model = CommentScorer()
    prediction = model.predict(features)
    return {"prediction": round(prediction, 2),
            "input": data.dict()}
```

Endpoint for textual input

```
@app.post("/analyze_text")  
  
def analyze(comment: CommentText):  
    forbidden = ["spam", "hate", "free"  
                "fake", "sign up"]  
    text_lower = comment.lower()  
    issues = [word for word in forbidden  
              if word in text_lower]  
  
    return {  
        "issues": issues,  
        "needs_moderation": len(issues)  
    }
```

Output for comment: Sign up for free

```
{  
    "issues": ["free", "sign up"],  
    "needs_moderation": 2  
}
```

Let's practice!

DEPLOYING AI INTO PRODUCTION WITH FASTAPI

Input validation in FastAPI

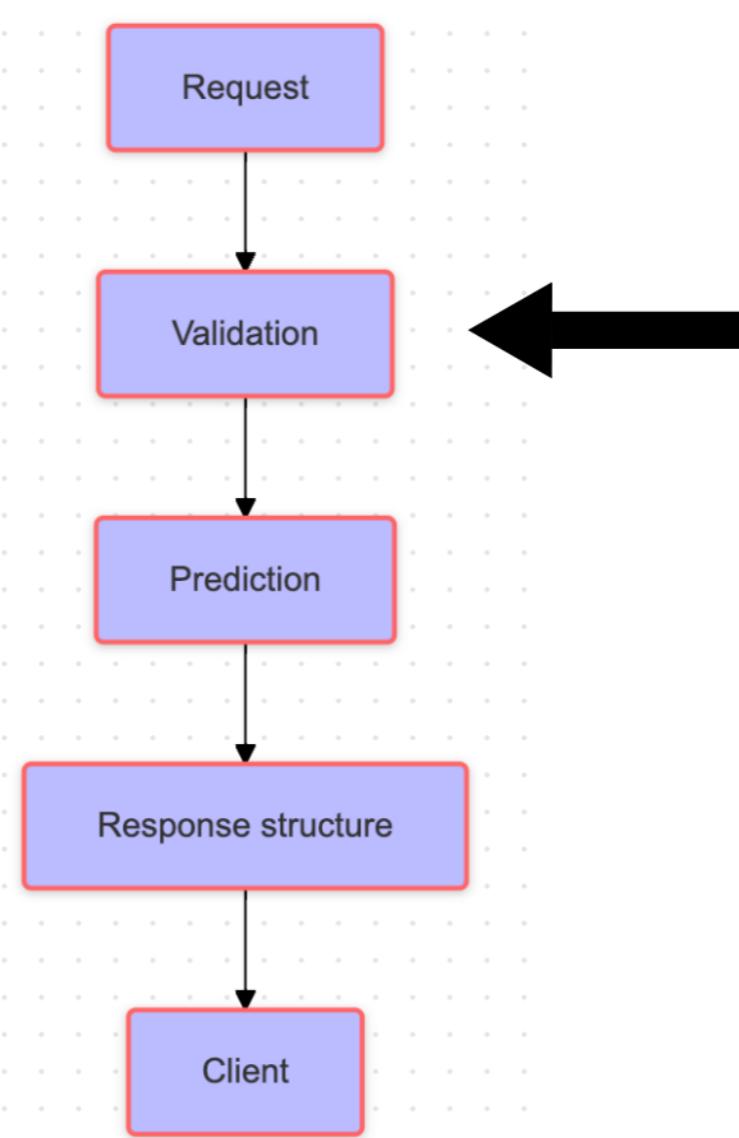
DEPLOYING AI INTO PRODUCTION WITH FASTAPI



Matt Eckerle

Software and Data Engineering Leader

Validating input data



Why validate the input?

- Validation for data integrity
- Prevent errors in the application
- Integrates with Pydantic
- Provided powerful tools for data validation



Pydantic for pre-defined function



Field Validators

The `Field` function is used to customize and add metadata to fields of models

Custom validation with pydantic



Field Validators

The `Field` function is used to customize and add metadata to fields of models



Custom Domain-Specific Validators

Create and apply custom validator functions

Graceful error reporting



Field Validators

The `Field` function is used to customize and add metadata to fields of models



Custom Domain-Specific Validators

Create and apply custom validator functions



Validation Error Handling

Custom messages and user-friendly reporting

Pydantic field validators

- User registration endpoint
- Validating the username entered by users:

```
from pydantic import BaseModel, Field
```

```
class User(BaseModel):  
    username: str = Field(..., min_length=3, max_length=50)
```

Adding custom validators

```
class User(BaseModel):  
    username: str = Field(...,  
                          min_length=3,  
                          max_length=50)  
  
    age: int  
  
    @field_validator('age')  
    def age_criteria(cls, age):  
        if age < 13:  
            raise ValueError('User must be at least 13')  
  
        return age
```

Custom validators in action

Valid request:

```
{"username": "john_doe", "age": 25}
```

```
Valid user: username='john_doe' age=25
```

Invalid request:

```
{"username": "too_young", "age": 10}
```

```
Validation error for {'username': 'too_young', 'age': 10}: User must be at least 13
```

Putting it all together



Field Validators

The `Field` function is used to customize and add metadata to fields of models



Custom Domain-Specific Validators

Create and apply custom validator functions



Validation Error Handling

Custom messages and user-friendly reporting

- Field validator for username
- Custom validator for age
- Error message if failing validation

Putting it all together

```
@app.post("/users")
def create_user(user: User):
    return {"message": "User created",
            "user": user.model_dump()}
```

Output:

```
{
    "message": "User created successfully",
    "user": {
        "username": "john_doe",
        "age": 25
    }
}
```

Let's practice!

DEPLOYING AI INTO PRODUCTION WITH FASTAPI

Loading a pre-trained model

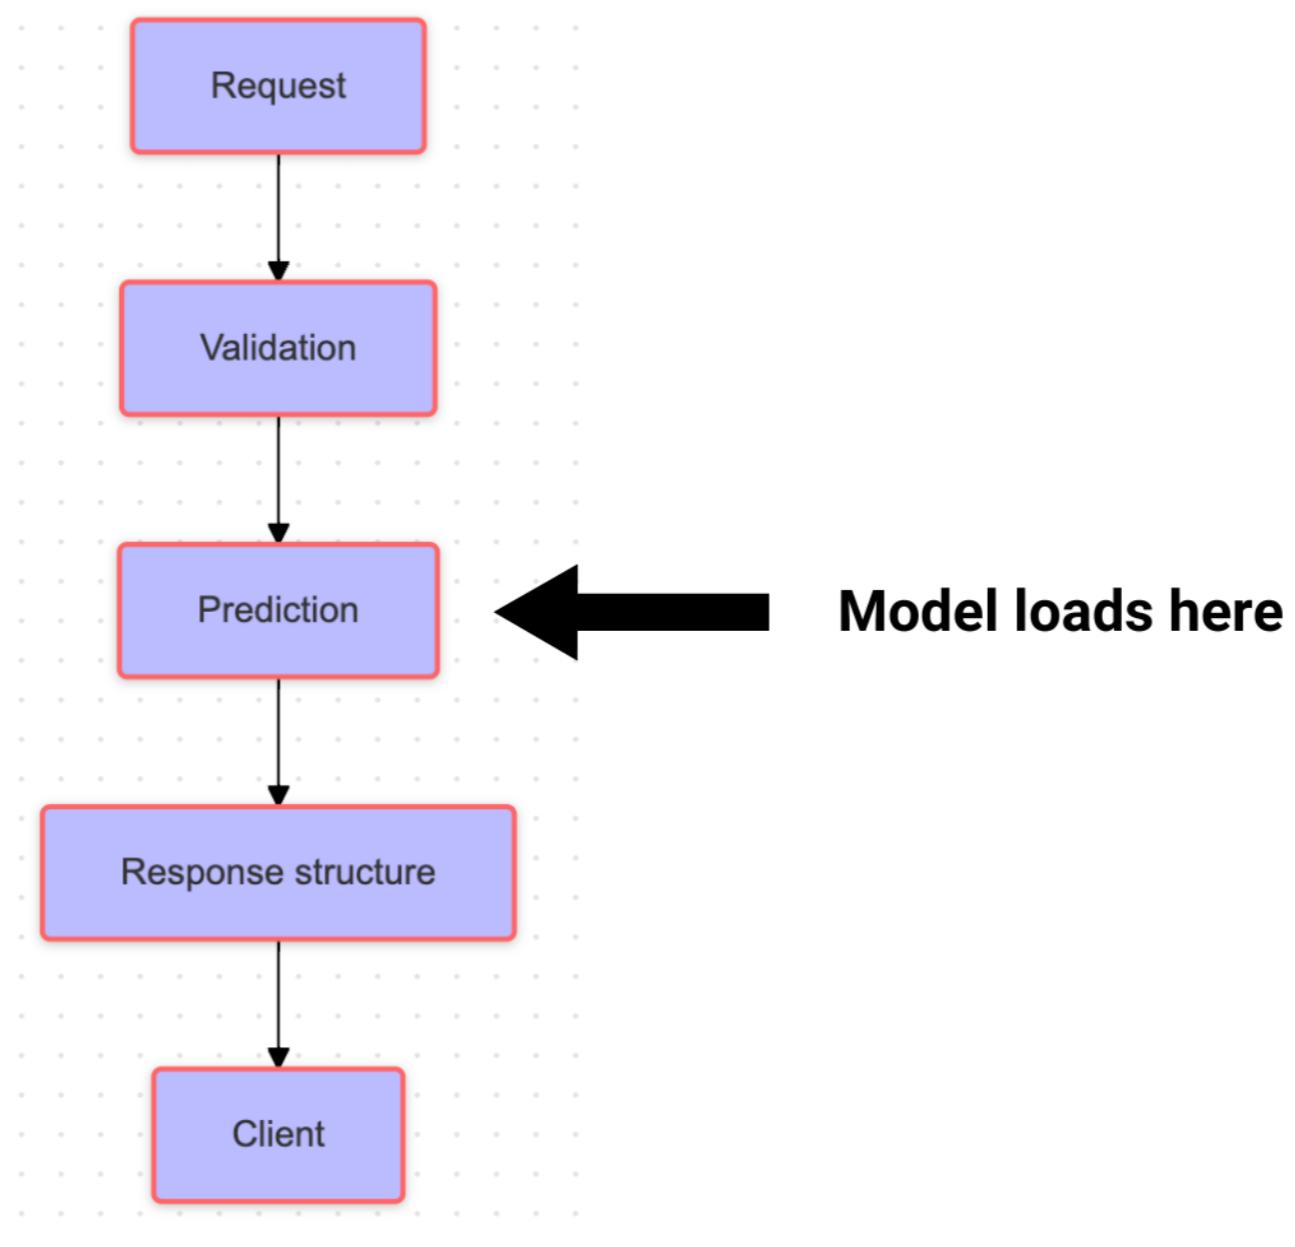
DEPLOYING AI INTO PRODUCTION WITH FASTAPI



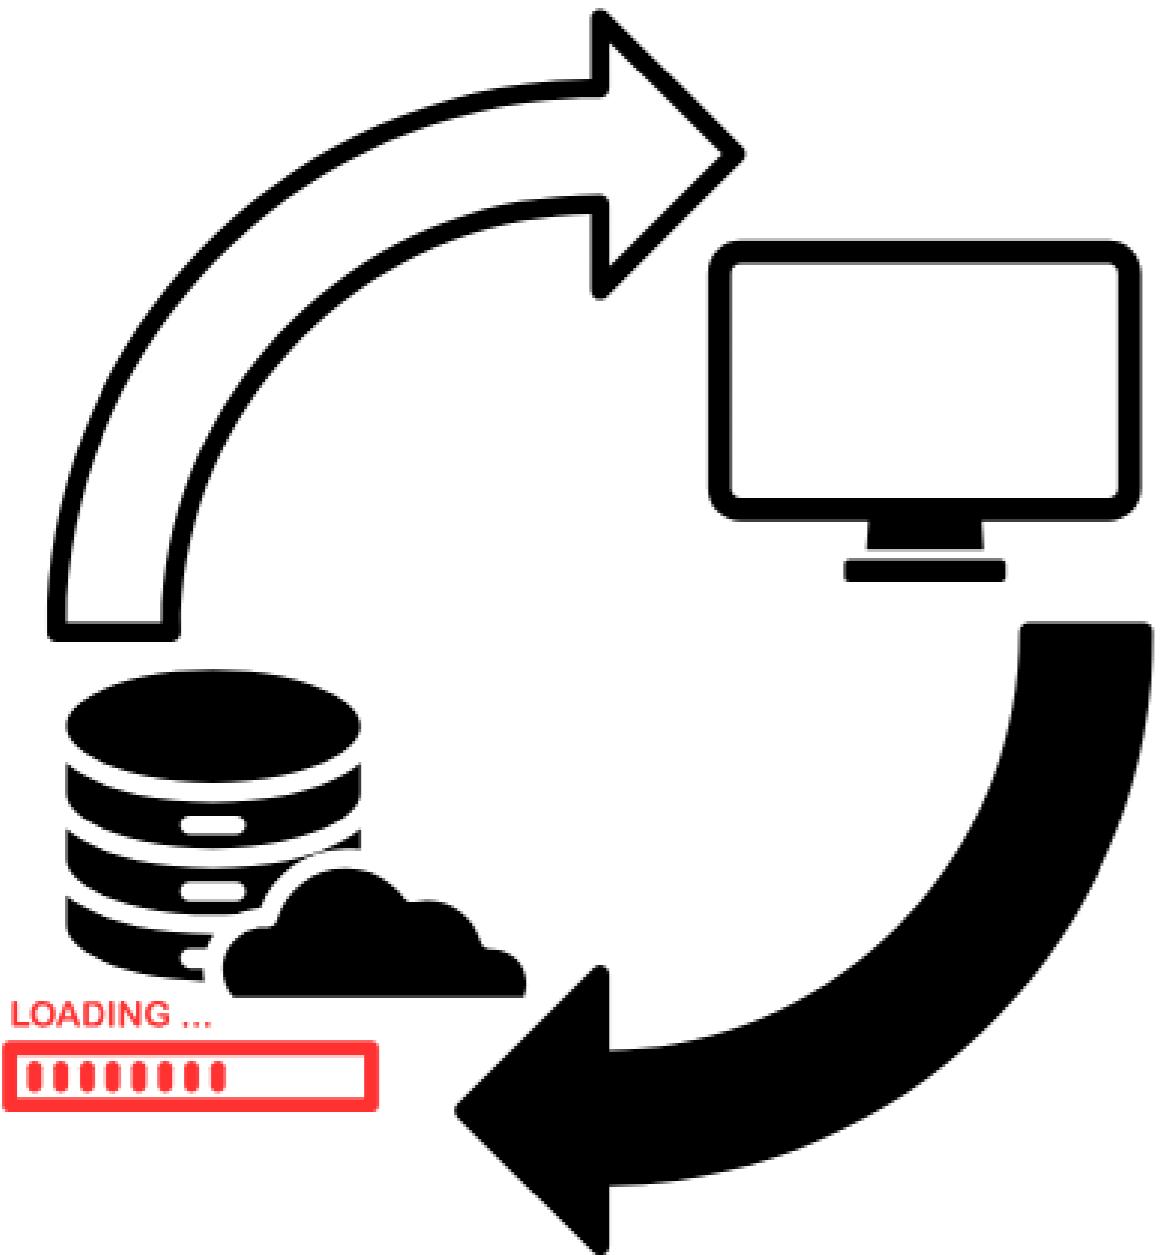
Matt Eckerle

Software and Data Engineering Leader

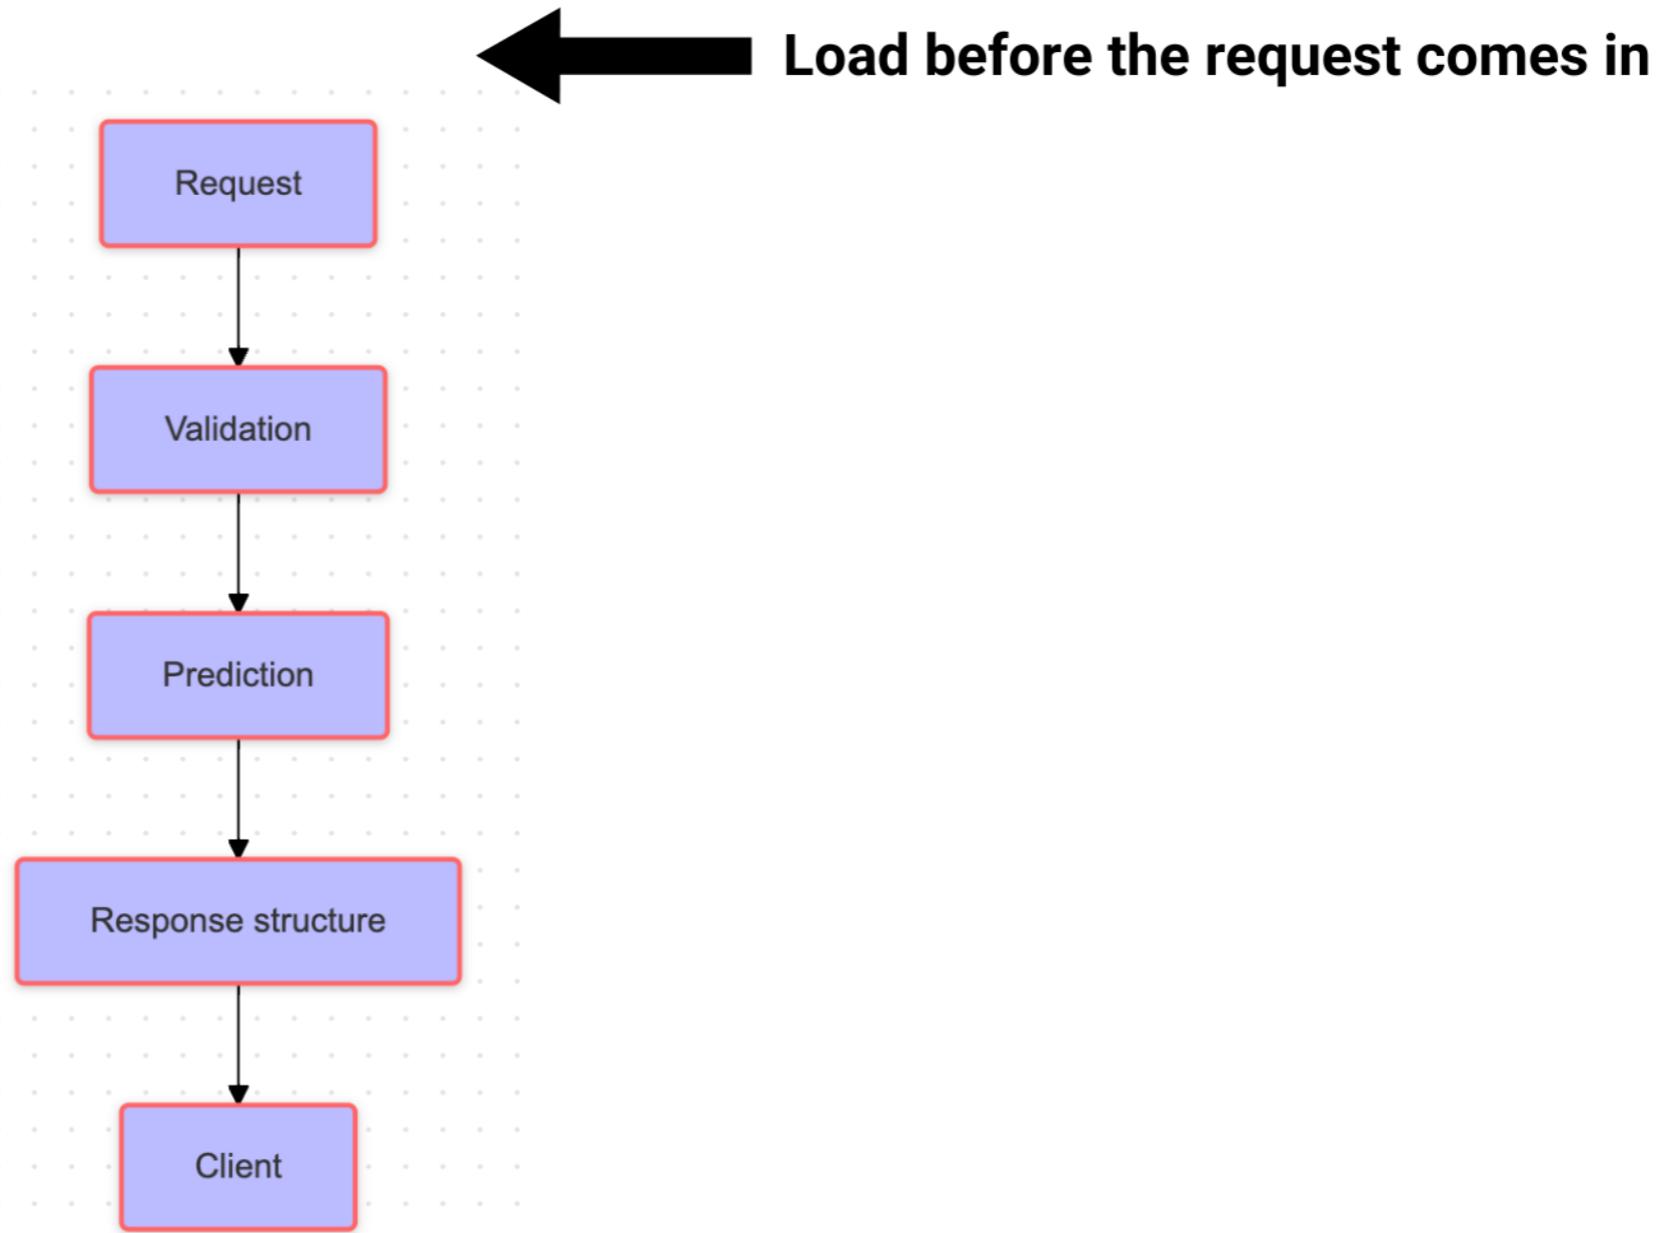
Current structure



Challenge with loading models



Load models before the request



Loading the model

```
from fastapi import FastAPI
```

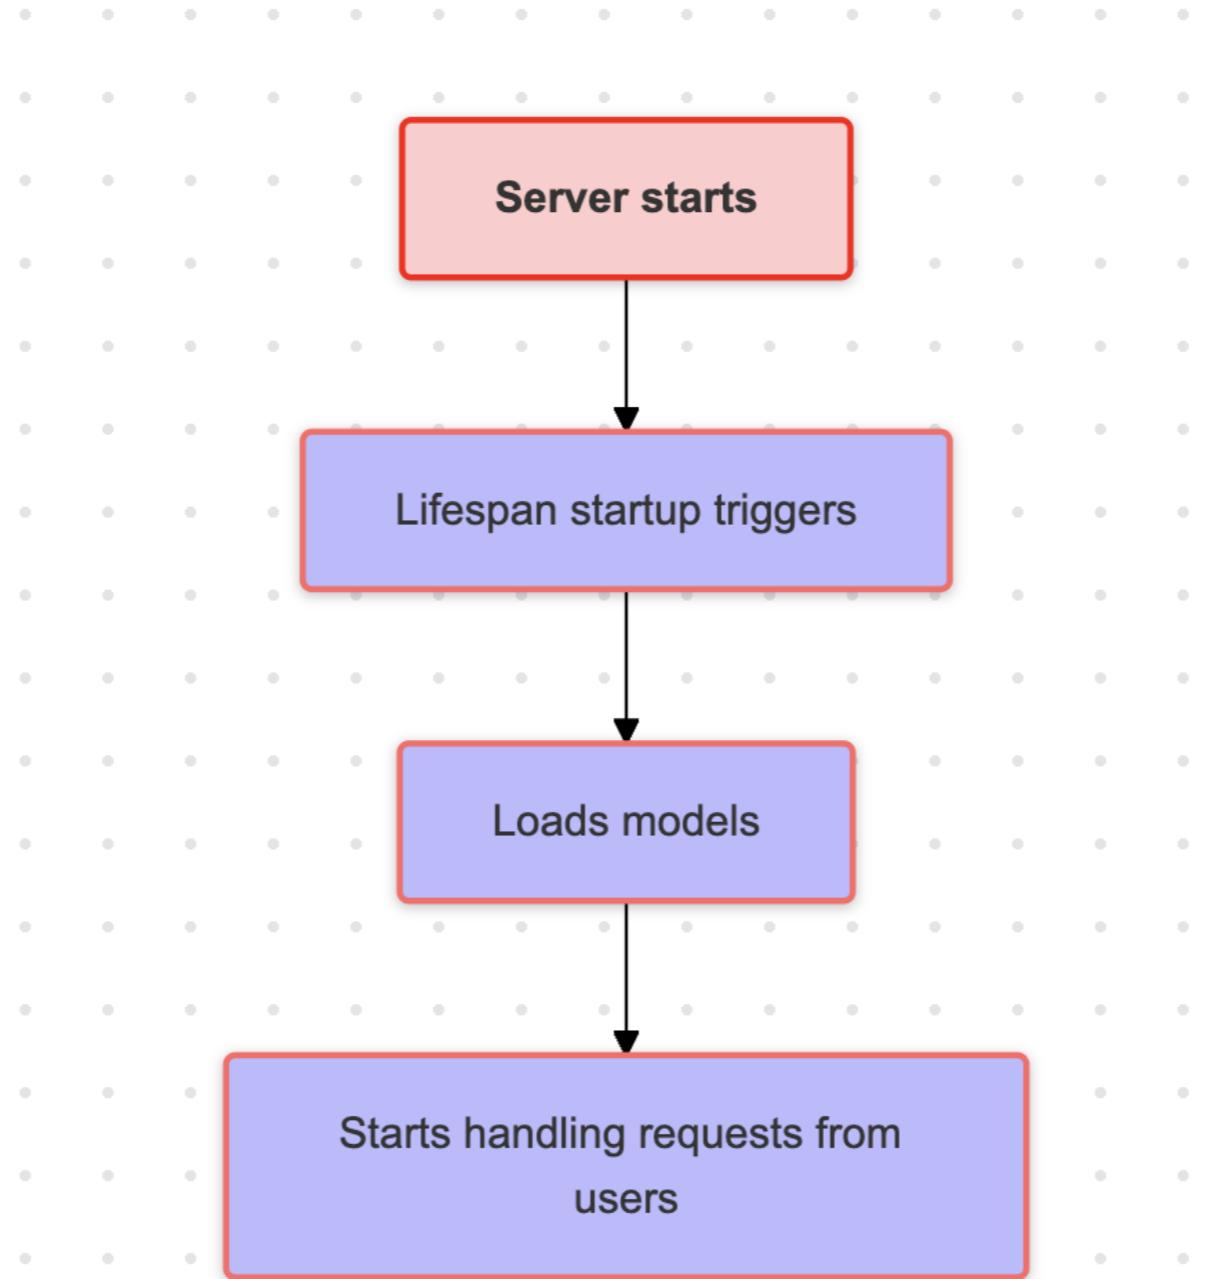
```
sentiment_model = None
```

```
def load_model():
    global sentiment_model
    sentiment_model = SentimentAnalyzer("trained_model.joblib")
    print("Model loaded successfully")
```

```
load_model()
```

```
Model loaded successfully
```

FastAPI lifespan event



FastAPI lifespan event

```
from contextlib import asynccontextmanager
```

```
@asynccontextmanager
async def lifespan(app: FastAPI):
    # Startup: Load the ML model
    load_model()
    yield
```

```
app = FastAPI(lifespan=lifespan)
```

Health checks

```
@app.get("/health")  
def health_check():  
    if sentiment_model is not None:  
        return {"status": "healthy",  
                "model_loaded": True}  
  
    return {"status": "unhealthy",  
            "model_loaded": False}
```

Curl command:

```
curl -X GET \  
"http://localhost:8080/health" \  
-H "accept: application/json"
```

Output:

```
{  
    "status": "healthy",  
    "model_loaded": true  
}
```

Let's practice!

DEPLOYING AI INTO PRODUCTION WITH FASTAPI

Returning structured prediction response

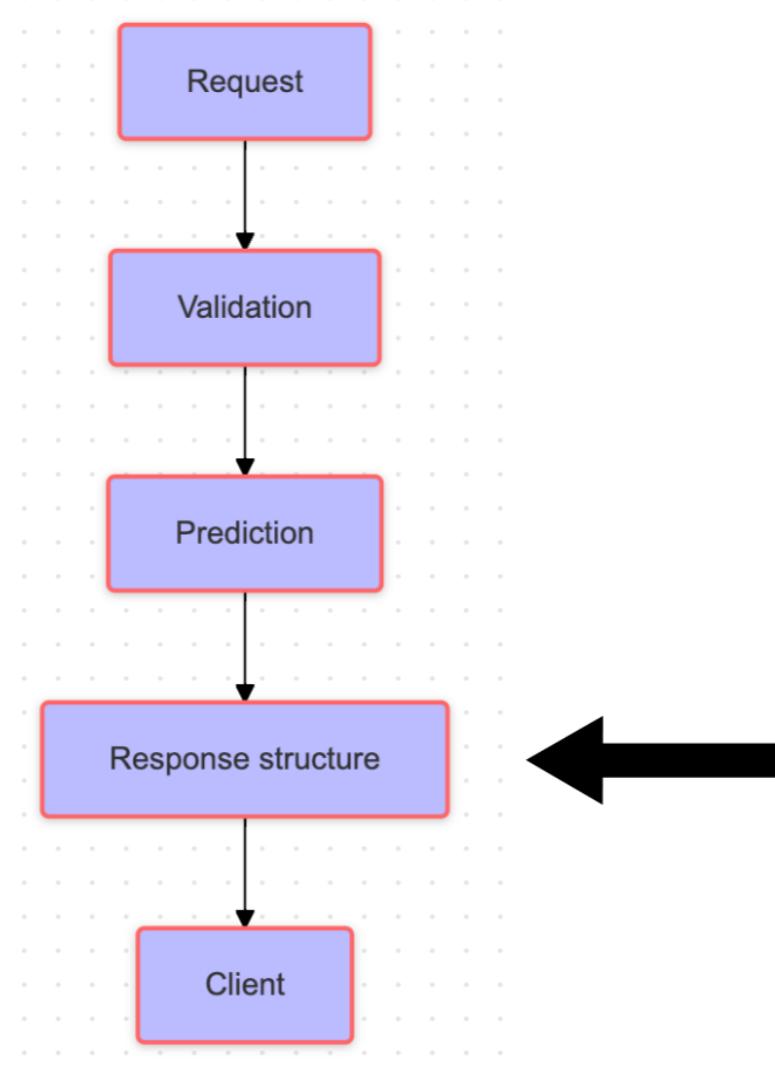
DEPLOYING AI INTO PRODUCTION WITH FASTAPI



Matt Eckerle

Software and Data Engineering Leader

Challenges with deploying models



1. Accept input data properly
2. Validate incoming data and handle errors
3. Make predictions
4. Return well-structured responses

Defining request structure

```
from pydantic import BaseModel
```

```
class PredictionRequest(BaseModel):  
    text: str
```

```
class PredictionResponse(BaseModel):  
    text: str  
    sentiment: str  
    confidence: float
```

Creating the prediction endpoint

```
@app.post("/predict")
def predict_sentiment(request: PredictionRequest):
    if sentiment_model is None:
        raise HTTPException(
            status_code=503,
            detail="Model not loaded"
        )
    result = sentiment_model(request.text)
    return PredictionResponse(
        text=request.text,
        sentiment=result[0]["label"],
        confidence=result[0]["score"]
    )
```

Input JSON:

```
{"text": "This movie was fantastic!"}
```

Response:

```
{
    "text": "This movie was fantastic!",
    "sentiment": "POSITIVE",
    "confidence": 0.95
}
```

Error handling

```
try:  
  
    result = sentiment_model(request.text)  
  
    return PredictionResponse(  
        text=request.text,  
        sentiment=result[0]["label"],  
        confidence=result[0]["score"]  
  
    )  
  
except Exception:  
  
    raise HTTPException(  
        status_code=500,  
        detail="Prediction failed"  
    )
```

Response when model fails to predict

```
{  
    "detail": "Prediction failed",  
    "status_code": 500  
}
```

Testing the endpoint

```
# Example request
import requests

response = requests.post(
    "http://localhost:8000/predict",
    json={"text": "Great product!"}
)
print(response.json())
```

```
{
  "text": "Great product!",
  "sentiment": "POSITIVE",
  "confidence": 0.998
}
```

Let's practice!

DEPLOYING AI INTO PRODUCTION WITH FASTAPI